

Contents

1	arm-fp-emu.c	2
2	assembly.h	8
3	fpuemu.h	18
4	rmaps.h	20
5	debug-print.h	23

1 arm-fp-emu.c

Called by the dynamic linker/loader and where the instrumentation process begins.

```
1 #include <assert.h>
2 #include <stdio.h>
3 #include <string.h>
4 #include <stdint.h>
5 #include <link.h>
6 #include "librunt.h"
7 #include "maps.h"
8 #include "dso-meta.h"
9 #include "relf.h"
10 #include "debug-print.h"
11 #include "rmaps.h"
12 #include "assembly.h"
13
14 // Fixes undefined symbols at build stage: the —defsym compiler flag
   doesn't solve this.
15 char* __private_strdup(const char *s) { return strdup(s); }
16 void* __private_malloc(size_t size) { return malloc(size); }
17
18 /*
19  * Sets the write flag for a region of memory.
20  * A 'guarantee' pointer ensures that this region includes that
21  * address.
22  */
23 int make_writable(void* from, void* to, void* guarantee) {
24
25     from = ROUND_DOWN_PTR_TO_PAGE(from);
26     if (guarantee != NULL && (guarantee < from || guarantee > to)) {
27         printfdbg("ERROR: Couldn't make region writable\n");
28         return -1;
29     }
30
31     size_t len = to - from;
32     int perms = PROT_READ | PROT_WRITE | PROT_EXEC;
33
34     printfdbg("mprotect(%p, %d, rwx)\n", from, len);
35     int ret = mprotect(from, len, perms);
36
37     if (ret != 0) {
38         printfdbg("ERROR: Couldn't make region writable %s\n");
```

```

39         perror("mprotect");
40         return -1;
41     }
42     return 0;
43 }
44
45 /*
46 * Returns a pointer to the end of an ELF header
47 */
48 void* end_of_header(void* sections_start) {
49     Elf32_Ehdr* header = (Elf32_Ehdr*) sections_start;
50
51     // Elf32_Off is typedef'd in elf.h as uint32_t:
52     // "typedef uint32_t Elf32_Off;"
53     uint32_t ph_table_offset = header->e_phoff;
54     // Elf32_Half is also uint16_t
55     uint16_t ph_table_len = ((uint16_t) header->e_phnum) * ((uint16_t)
        header->e_phentsize);
56
57     return ((int8_t*) sections_start) + ph_table_offset + ph_table_len;
58 }
59
60 /*
61 * Attempts to set the write flag for a region of memory.
62 * Returns whether change was successful (0 == success)
63 */
64 int try_set_mem_writable(struct maps_entry *maps_ent, void* seg_start,
    void* seg_end) {
65     if (maps_ent->w != 'w') {
66         int ret = make_writable(seg_start, seg_end, NULL);
67         if (ret == 0) {
68             printfdbg("Write perms now set for %p-%p.\n",
                ROUND_DOWN_PTR_TO_PAGE(seg_start), seg_end);
69             return 0;
70         } else {
71             printfdbg("ERROR: Failure to make writable\n");
72             exit(1);
73         }
74     } else {
75         printfdbg("Write perms already set (%c%c%c), continuing.\n",
            maps_ent->r, maps_ent->w, maps_ent->x);
76         return 1;
77     }
78 }
79
80 /*

```

```

81 * Core of the instrumentation process.
82 * Looks through a mapped region in memory and finds a floating-point
    instruction.
83 * If one is found, it is replaced by a branch instruction and a
    trampoline is made and
84 * written to somewhere in memory.
85 */
86 void replace_instrs_in_segment(struct maps_entry *maps_ent, void* from,
    void* to) {
87     void* sections_start = from;
88     void* sections_end = to;
89     void* seg_start = maps_ent->first;
90     void* seg_end = maps_ent->second;
91
92     printfdbg("\tWe have entered replace_instructions() \n");
93     printfdbg("\tRange of executable instructions (inside segment range
        ): \n");
94     printfdbg("\t");
95     if (seg_start != sections_start) {
96         printfdbg("(%p-", seg_start);
97     }
98     printfdbg("%p-%p", sections_start, sections_end);
99     if (seg_end != sections_end) {
100         printfdbg("(-%p)", seg_end);
101     }
102     printfdbg("\n");
103
104     void* instrs_start = sections_start;
105     if (0 == strncmp(sections_start, "\177ELF", 4)) {
106         // Probably an ELF header so skip it
107         instrs_start = end_of_header(sections_start);
108         printfdbg("ELF headed spotted at %p, skipping to %p\n",
            sections_start, instrs_start);
109     }
110
111     assert(instrs_start >= seg_start && instrs_start <= seg_end);
112     assert(sections_end >= seg_start && sections_end <= seg_end);
113
114     int b_is_writable = maps_ent->w == "w";
115     int b_did_perm_change = 0;
116
117     printfdbg("Scanning through %p-%p for FP instructions\n",
        instrs_start, sections_end);
118     for (int8_t* instr = instrs_start; instr < sections_end-4; instr +=
        2) {
119         void* tramp = generate_trampoline(instr);

```

```

120         if (tramp == NULL) {
121             continue;
122         } else {
123             printfdbg("Trampoline written for FP instruction at %p in %
124                 p-%p\n", instr, instrs_start, sections_end);
125             printfdbg(" - writing jump at the mentioned instr (%p)\n",
126                 instr);
127             if (!b_is_writable && 0 == try_set_mem_writable(maps_ent,
128                 seg_start, seg_end)) {
129                 b_is_writable = 1;
130                 b_did_perm_change = 1;
131             }
132             insert_probe(instr, tramp);
133         }
134     }
135 }
136
137 /*
138  * Callback for librunt.
139  * Librunt passes information about a mapped region of memory and
140  * this callback method does some setup then instruments it.
141  *
142  * In full transparency, this method resembles code belonging to my
143  * supervisor Dr. Stephen Kell.
144  * His project "libsystap" uses his other project "librunt" which is
145  * also a dependency for this project.
146  * Not only does some of the code look similar because they are using
147  * librunt in mostly the same way because
148  * they both are rewriting code in memory, but libsystap contributed to
149  * what I used to learn about instrumentation
150  * and how to use librunt.
151  * The relevant code is spread throughout the file "src/trap.c" which
152  * can be found at the link below. This includes the methods
153  * "trap_one_executable_region_given_shdrs", "trap_one_executable_region
154  * ", and "trap_one_instruction_range".
155  * https://github.com/stephenrkell/libsysstrap/blob/21
156  \* c5b00eb256f5489ee0d163efecc3398dfef2c9/src/trap.c
157  */
158 static int handle_maps_entry(struct maps_entry *maps_ent, char *linebuf
159     , void *arg) {
160     // Tests for memory regions that don't need instrumenting.
161     // "[" is here to ensure stability by exercising overcaution, but
162     // in practice you'd want
163     // to be more specific like the rest of the search strings.
164     char* to_skip[] = {"[", "[stack]", "[vvar]", "[sigpage]", "[vdso]",
165         "[vectors]", "libm-2.31.so", "libkeystone.so.0", "libcapstone.

```

```

153         so.4" };
154     for (int i = 0; i < sizeof(to_skip) / sizeof(to_skip[0]); i++) {
155         if (NULL != strstr(maps_ent->rest, to_skip[i])) {
156             printfdbg("\tSkipping %s\n", maps_ent->rest);
157             return 0;
158         }
159     }
160     printfdbg("Handling maps entry for \"%s\"\n", maps_ent->rest);
161     printfdbg("\tGetting file metadata.\n");
162     struct file_metadata* meta = __runt_files_metadata_by_addr(maps_ent
163     ->first);
164     if (meta == NULL) {
165         printfdbg("File metadata not found for: %s\n", maps_ent->rest);
166         return 1;
167     } else if (maps_ent->x != 'x') {
168         printfdbg("\tNot executable, skipping.\n");
169         return 0;
170     }
171     // "Base address shared object is loaded at" - definition
172     ElfW(Addr) l_addr = meta->l->l_addr;
173     if (meta->shdrs == NULL) {
174         printfdbg("\tNo section headers: file-metadata->shdrs is null.\n");
175         return 1;
176     } else {
177         printfdbg("\tFound section headers.\n");
178     }
179     // Range within segment containing exactly only sections
180     void* sections_from = l_addr + find_section_boundary((uintptr_t) (
181     maps_ent->first - l_addr), \
182     SHF_EXECINSTR, \
183     0, \
184     meta->shdrs, \
185     meta->ehdr->e_shnum, \
186     NULL);
187     void* sections_to = l_addr + find_section_boundary((uintptr_t) (
188     maps_ent->second - l_addr), \
189     SHF_EXECINSTR, \
190     1, \
191     meta->shdrs, \
192     meta->ehdr->e_shnum, \
193     NULL);

```

```

194     printfdbg("%s\n", maps_ent->rest);
195     assert(maps_ent->first <= sections_from && sections_from <=
           sections_to && sections_to <= maps_ent->second);
196     printfdbg("Maps entry goes from (%p-%p-%p(-%p))\n", maps_ent->first
           , sections_from , sections_to , maps_ent->second);
197
198     replace_instrs_in_segment(maps_ent, sections_from , sections_to);
199     return 0;
200 }
201
202 /*
203  * Called by the dynamic linker/loader before the base program.
204  * This is where the instrumentation happens.
205  */
206 static int entrypoint(void) __attribute__((constructor(101)));
207 static int entrypoint(void) {
208     int fd = open_maps();
209     start_disasm_engine();
210     start_asm_engine();
211     emulator_init();
212
213     // Replace instructions
214     process_all_lines(fd , handle_maps_entry);
215
216     close_maps();
217     stop_disasm_engine();
218     stop_asm_engine();
219     return 0;
220 }

```

2 assembly.h

Where assembly and disassembly takes place as well as methods to write and link up trampolines.

```
1 #include <keystone/keystone.h>
2 #include <capstone/capstone.h>
3 #include <inttypes.h>
4 #include <sys/mman.h>
5 #include <stddef.h>
6 #include "fpuemu.h"
7 #include <unistd.h>
8 #include <stdlib.h>
9 #define PAGE_SIZE sysconf(_SC_PAGE_SIZE)
10
11 int8_t MIN_PRINTABLE_ASCII = (int8_t) 0x20; // 32
12 int8_t MAX_PRINTABLE_ASCII = (int8_t) 0x7F; // 127
13 int INT24_MIN = -(1 << 23);
14 int INT24_MAX = (1 << 23) - 1;
15
16 /*
17  * Template for the trampolines used to connect
18  * a probe site and emulation routine at run-time.
19  */
20 int8_t template[] = {
21     0xFF, 0x5F, 0x2D, 0xE9, // push {r0-r12, r14}
22     0xAD, 0x5E, 0x0D, 0xE3, // movw r5, #0xdead
23     0x05, 0x58, 0xA0, 0xE1, // lsl r5, r5, #16
24     0xEF, 0x6E, 0x0B, 0xE3, // movw r6, #0xbeef
25     0x06, 0x50, 0x85, 0xE1, // orr r5, r5, r6
26     0x00, 0x00, 0xA0, 0xE3, // mov r0, #0
27     0x00, 0x10, 0xA0, 0xE3, // mov r1, #0
28     0x00, 0x20, 0xA0, 0xE3, // mov r2, #0
29     0x00, 0x30, 0xA0, 0xE3, // mov r3, #0
30     0x35, 0xFF, 0x2F, 0xE1, // blx r5
31     0xFF, 0x5F, 0xBD, 0xE8, // pop {r0-r12, r14}
32     0xFE, 0xFF, 0xFF, 0xEA // b #0
33 };
34
35 /*
36  * Offsets of various instructions in the trampoline
37  * template and the names of registers used.
38  */
39 int MOV_UPPER_OFFSET = 1 * 4;
```



```

40 int MOV_LOWER_OFFSET = 3 * 4;
41 int MOV_R0_OFFSET = 5 * 4;
42 int MOV_R1_OFFSET = 6 * 4;
43 int MOV_R2_OFFSET = 7 * 4;
44 int MOV_R3_OFFSET = 8 * 4;
45 int RET_OFFSET = 11 * 4;
46 int REG_SCRATCH = 6;
47 int REG_CALL = 5;
48
49 /*
50  * Capstone (disassembly framework) and
51  * Keystone (assembly framework) engine handles.
52  */
53 csh* cs_handle;
54 ks_engine* ks_handle;
55
56 /*
57  * Copy four bytes from one location into another.
58  * Used to replace four-byte ARM instructions.
59  */
60 void clobber(void* dst, void* src) {
61     printfdbg("clobbering\n");
62     make_writable(dst, dst + 4, NULL);
63     printfdbg(" - writing src into dst\n");
64     memcpy(dst, src, 4);
65 }
66
67 void start_asm_engine() {
68     if(ks_open(KS_ARCH_ARM, KS_MODE_ARM, &ks_handle) != KS_ERR_OK) {
69         printfdbg("\tUnable start the keystone assembly engine.\n");
70     }
71 }
72
73 void stop_asm_engine() {
74     ks_close(ks_handle);
75 }
76
77 void start_disasm_engine() {
78     if(cs_open(CS_ARCH_ARM, CS_MODE_ARM, &cs_handle) != CS_ERR_OK) {
79         printfdbg("\tUnable start the capstone disassembly engine.\n");
80         exit(-1);
81     }
82     cs_malloc(cs_handle);
83
84     // enable full range of disassembly information
85     cs_option(cs_handle, CS_OPT_DETAIL, CS_OPT_ON);

```

```

86 }
87
88 void stop_disasm_engine() {
89     cs_close(cs_handle);
90 }
91
92 /*
93  * Use the Keystone assembler to convert a string in
94  * assembly to machine-code.
95  */
96 int8_t* assemble_instr(char* assembly) {
97     int8_t* instr;
98     size_t size;
99     size_t count;
100     if (ks_asm(ks_handle, assembly, 0, &instr, &size, &count) !=
        KS_ERR_OK) {
101         printfdbg("Unable to assemble instruction '%s'\n", assembly);
102         exit(-1);
103     }
104     return instr;
105 }
106
107 /*
108  * Returns a pointer to a machine-code branch
109  * instruction (ARM) that branches to the given offset.
110  * 'offset' is a human-readable string like "0xfae".
111  */
112 int8_t* assemble_branch(char* offset) {
113     char assembly[100];
114     sprintf(assembly, "b #0x%s", offset);
115     return assemble_instr(assembly);
116 }
117
118 /*
119  * Returns a pointer to a machine-code 'mov'
120  * instruction (ARM). E.g. reg = 5 for r5. 'val' is
121  * the immediate value.
122  */
123 int8_t* assemble_mov(uint8_t reg, uint16_t val) {
124     int8_t instr[] = {
125         0xEF, 0x1E, 0x0B, 0xE3 // movw r1, #0xbeef
126     };
127
128     if (reg > 0xF) {
129         printfdbg("assemble_mov() - Register not supported: %lu", reg);
130         return NULL;

```

```

131     }
132
133     // set register
134     instr[1] = (reg & 0x0F) << 4;
135
136     // the immediate value is divided and placed in interesting
137     // positions
138     // (value's nibbles 0,1,2,3 to mov's nibbles 5,3,0,1)
139     // so various logical operations are needed
140     // Also it is a "byte" array not a nibble array so OR-ing is needed
141     instr[2] = (instr[2] & 0xF0) | ((val & 0xF000) >> 12);
142     instr[1] = (instr[1] & 0xF0) | ((val & 0x0F00) >> 8);
143     instr[0] = val & 0x00FF;
144
145     int8_t* outbuf = malloc(4*sizeof(int8_t));
146     memcpy(outbuf, instr, sizeof(instr));
147     return outbuf;
148 }
149
150 /*
151 * Returns a struct that contains information about what instruction
152 * exists at the provided pointer.
153 */
154 cs_insn* disassemble_instr(void* p_instr) {
155     int8_t code[4];
156     memcpy(code, (int8_t*) p_instr, 4);
157
158     cs_insn* instr;
159     assert(cs_handle != NULL);
160     int count = cs_disasm(cs_handle, p_instr, 4, 0, 1, &instr);
161
162     cs_insn* output;
163     if (count == 0 || 0 == strcmp(instr[0].mnemonic, "")) {
164         output = NULL;
165     } else {
166         cs_insn* insn_copy = malloc(sizeof(cs_insn));
167         memcpy(insn_copy, &instr[0], sizeof(cs_insn));
168         output = insn_copy;
169         cs_free(instr, count);
170     }
171
172     char* name = output == NULL ? NULL : cs_insn_name(cs_handle, output
173     ->id);
174     return output;
175 }

```

```

175  /*
176  * Returns the name of the instruction at the
177  * provided pointer.
178  */
179  char* instr_name(void* instr){
180      cs_insn* disassembly = disassemble_instr(instr);
181      if (disassembly == NULL) {
182          return NULL;
183      }
184      char* name = cs_insn_name(cs_handle, disassembly->id);
185      char* out = malloc(10);
186      strcpy(out, name);
187      free(disassembly);
188      return out;
189  }
190
191  /*
192  * In the instrumentation stage, this method displaces the floating-
193  * point
194  * instruction with a branch that points to the start of a pre-written
195  * trampoline.
196  * The trampoline is also written to so that when it returns, it returns
197  * to just after
198  * the displaced FP instruction.
199  */
200  int insert_probe(void* instr, void* tramp) {
201      printfdbg("Inserting probe at %p to connect to trampoline at %p\n",
202          instr, tramp);
203
204      // Calculate trampoline offset
205      ptrdiff_t offset = ((int8_t*) tramp - (int8_t*) instr);
206      ptrdiff_t offset_reverse = (((int8_t*) instr + 4) - ((int8_t*)
207          tramp + RET.OFFSET));
208      // Ensure trampoline is close enough for the offset to be written
209      assert(INT24_MIN <= offset && offset <= INT24_MAX);
210      assert(INT24_MIN <= offset_reverse && offset_reverse <= INT24_MAX);
211
212      // Convert offset into a string to pass to the assembler
213      char str_offset[100];
214      char str_offset_reverse[100];
215      sprintf(str_offset, "%td", offset);
216      sprintf(str_offset_reverse, "%td", offset_reverse);
217      int8_t* probe_site_to_trampoline = assemble_branch(str_offset);
218      int8_t* tramp_to_probe_site = assemble_branch(str_offset_reverse);
219
220      #ifdef DO_DBG_PRINT

```

```

216     char* before = instr_name(instr);
217     char* after = instr_name(probe_site_to_trampoline);
218     char* back_again = instr_name(trampoline_to_probe_site);
219
220     printfdbg(" - writing tramp branch into instr at %p\n", instr);
221     printfdbg("      - instr: %02x%02x%02x%02x",
222             *((int8_t*) instr)&0xff,
223             *((int8_t*) instr + 1) & 0xff,
224             *((int8_t*) instr + 2) & 0xff,
225             *((int8_t*) instr + 3) & 0xff);
226     printfdbg(" (%s)\n", before);
227     printfdbg("      - assembly: %02x%02x%02x%02x",
228             *((int8_t*) probe_site_to_trampoline)&0xff,
229             *((int8_t*) probe_site_to_trampoline + 1) & 0xff,
230             *((int8_t*) probe_site_to_trampoline + 2) & 0xff,
231             *((int8_t*) probe_site_to_trampoline + 3) & 0xff);
232     printfdbg(" (%s)\n", after);
233
234     printfdbg(" - writing return branch into tramp at %p\n", (int8_t*)
235             tramp + RET_OFFSET);
236     printfdbg("      - instr: %02x%02x%02x%02x",
237             *((int8_t*) tramp + RET_OFFSET)&0xff,
238             *((int8_t*) tramp + RET_OFFSET + 1) & 0xff,
239             *((int8_t*) tramp + RET_OFFSET + 2) & 0xff,
240             *((int8_t*) tramp + RET_OFFSET + 3) & 0xff);
241     printfdbg(" (%s)\n", before);
242     printfdbg("      - assembly: %02x%02x%02x%02x",
243             *((int8_t*) probe_site_to_trampoline) & 0xff,
244             *((int8_t*) probe_site_to_trampoline + 1) & 0xff,
245             *((int8_t*) probe_site_to_trampoline + 2) & 0xff,
246             *((int8_t*) probe_site_to_trampoline + 3) & 0xff);
247     printfdbg(" (%s)\n", after);
248     #endif
249
250     // Replace FP instruction with branch
251     clobber(instr, probe_site_to_trampoline);
252
253     // Add return branch to the end of the trampoline
254     clobber((int8_t*)trampoline + RET_OFFSET, trampoline_to_probe_site);
255
256     #ifdef DO_DBG_PRINT
257     printfdbg(" - branch written\n");
258     printfdbg(" - Therefore, '%s' replaced with '%s' at %p\n", before,
259             after, instr);
260     free(before);
261     free(after);

```

```

260     free(back_again);
261     #endif
262
263     ks_free(probe_site_to_trampoline);
264     ks_free(trampoline_to_probe_site);
265 }
266
267 /*
268  * Finds and reserves a page of memory near 'instr_addr'.
269  * Returns a pointer to the start of this page.
270  */
271 void* mmap_nearby(void* instr_addr) {
272     unsigned int perms = PROT_EXEC | PROT_READ | PROT_WRITE;
273     unsigned int flags = MAP_FIXED_NOREPLACE | MAP_PRIVATE |
274         MAP_ANONYMOUS;
275
276     void* range_low = 0x0;
277     void* range_high = 0xFFFFFFFF;
278     if (instr_addr >= -INT24_MIN) {
279         range_low = instr_addr + INT24_MIN;
280     }
281     if (instr_addr <= range_high - INT24_MAX) {
282         range_high = instr_addr + INT24_MAX;
283     }
284     assert(range_low <= range_high);
285     void* search_from = ROUND_UP_PTR_TO_PAGE(range_low);
286     void* search_to = ROUND_DOWN_PTR_TO_PAGE(range_high);
287     printfdbg("mmap_nearby: We have %p-%p range but will look at pages
288         %p-%p\n", range_low, range_high, search_from, search_to);
289
290     assert(search_from <= instr_addr && instr_addr <= search_to);
291
292     void* map_region = NULL;
293     printfdbg("Searching from %p to %p\n", search_from, search_to);
294     for (int page_start = search_from; page_start < search_to;
295         page_start += PAGE_SIZE) {
296         printfdbg("mmap(%p, %d, ...) = ", page_start, sizeof(template))
297             ;
298         map_region = mmap(page_start, sizeof(template), perms, flags,
299             -1, 0);
300         printfdbg("%p (should be %p)\n", map_region, page_start);
301         if (map_region == page_start) { // request accepted
302             break;
303         }
304     }
305     #ifdef DO_DBG_PRINT
306     perror("mmap");
307     #endif

```

```

301         #endif
302     }
303     if (map_region == MAP_FAILED || map_region == NULL) {
304         printfdbg("ERROR: no space for trampoline near instruction %p (
305             see mmap error below)\n", instr_addr);
306         #ifdef DO_DBG_PRINT
307             perror("mmap");
308         #endif
309         return NULL;
310     }
311     return map_region;
312 }
313 /*
314  * Reserves a page of memory and writes the trampoline template to it.
315  */
316 void* gen_template_trampoline(void* instr_addr) {
317     void* p_template = mmap_nearby(instr_addr);
318     if (p_template == NULL) {
319         printfdbg("ERROR: no space for trampoline near instruction %p (
320             see mmap error below)\n", instr_addr);
321         #ifdef DO_DBG_PRINT
322             perror("mmap");
323         #endif
324         return NULL;
325     }
326     memcpy(p_template, &template, sizeof(template));
327     return p_template;
328 }
329 /*
330  * Shorthand to replace an instruction in a trampoline.
331  */
332 void insert_trampoline_instr(void* trampoline, void* instr, int offset) {
333     memcpy(((int8_t*)trampoline) + offset, (int8_t*)instr, 4*sizeof(
334         int8_t));
335 }
336 /*
337  * Inserts the emulation routine arguments into a trampoline.
338  */
339 void trampoline_insert_emu_args(void* trampoline, int arg1, int arg2, int
340     arg3, int arg4) {
341     int8_t* mov_r0 = assemble_mov(0, arg1);
342     int8_t* mov_r1 = assemble_mov(1, arg2);
343     int8_t* mov_r2 = assemble_mov(2, arg3);

```

```

343     int8_t* mov_r3 = assemble_mov(3, arg4);
344     insert_tramp_instr(trampoline, mov_r0, MOV_R0-OFFSET);
345     insert_tramp_instr(trampoline, mov_r1, MOV_R1-OFFSET);
346     insert_tramp_instr(trampoline, mov_r2, MOV_R2-OFFSET);
347     insert_tramp_instr(trampoline, mov_r3, MOV_R3-OFFSET);
348     free(mov_r0);
349     free(mov_r1);
350     free(mov_r2);
351     free(mov_r3);
352 }
353
354 /*
355  * Writes the address of the emulation routine into the trampoline for
356  * branching later. As described in the report, the address is divided
357  * and
358  * loaded in two parts into a register by instructions in the trampoline
359  */
359 void link_tramp_to_emu(void* trampoline, void* func_address) {
360     // Split address into two two-byte pieces.
361     int16_t word_lower = ((int32_t) func_address) & 0x0000FFFF;
362     int16_t word_upper = (((int32_t) func_address) & 0xFFFF0000) >> 16;
363
364     int8_t* mov_instr0 = assemble_mov(REG_CALL, word_upper);
365     int8_t* mov_instr1 = assemble_mov(REG_SCRATCH, word_lower);
366
367     if (mov_instr0 == NULL || mov_instr1 == NULL) {
368         printfdbg("ERROR: couldn't assemble func address into mov");
369         exit(1);
370     }
371
372     insert_tramp_instr(trampoline, mov_instr0, MOV_UPPER-OFFSET);
373     insert_tramp_instr(trampoline, mov_instr1, MOV_LOWER-OFFSET);
374
375     free(mov_instr0);
376     free(mov_instr1);
377 }
378
379 /*
380  * Returns a pointer to the beginning of the trampoline or NULL if
381  * failed.
382  * In three places there is a hard-coded check for the 'vadd.f32'
383  * instruction as
384  * it is the only instruction emulated so far in the emulator. In a full
385  * solution,
386  * these checks would be removed.

```



```

384 */
385 void* generate_trampoline(void* instr_addr) {
386     cs_insn* disassembly = disassemble_instr(instr_addr);
387
388     if (disassembly == NULL) {
389         return NULL;
390     } else if (disassembly->id != ARM_INS_VADD) {
391         free(disassembly);
392         return NULL;
393     }
394
395     // Get which S registers are used
396     cs_arm* arm = &(disassembly->detail->arm);
397     int Sd = arm->operands[0].reg;
398     int Sn = arm->operands[1].reg;
399     int Sm = arm->operands[2].reg;
400
401     if (Sd != ARM_REG_S0 || Sn != ARM_REG_S0 || Sm != ARM_REG_S1 || arm
        ->cc != ARM_CC_AL) return NULL;
402     printfdbg("%s %s\n", disassembly->mnemonic, disassembly->op_str);
403     printfdbg("This vadd (CC=%d) instruction uses the registers %d, %d,
        %d %d %d\n", arm->cc, Sd, Sn, Sm, arm->operands[3].reg, arm->
        operands[4].reg);
404     assert(ARM_REG_S0 <= Sd && Sd <= ARM_REG_S31);
405     assert(ARM_REG_S0 <= Sn && Sn <= ARM_REG_S31);
406     assert(ARM_REG_S0 <= Sm && Sm <= ARM_REG_S31);
407
408     // Make trampoline
409     int8_t* tramp = gen_template_trampoline(instr_addr);
410     if (tramp == NULL) {
411         printfdbg("ERROR: failed to generate template trampoline\n");
412         exit(1);
413     }
414
415     // 'vadd.f32' is connected here as it is the only emulation routine
        present
416     // but in practice you'd want to check the disassembly above - the
        variable 'disassembly'.
417     link_trampoline_to_emu(tramp, &vadd_f32);
418
419     // Put the args (names of S registers) into r0-r2
420     tramp_insert_emu_args(tramp, Sd, Sn, Sm, 0);
421
422     printfdbg("Trampoline made for instruction.");
423     return tramp;
424 }

```

3 fpuemu.h

A stub floating-point emulator which implements virtual registers and an emulation routine for the ‘vadd.f32’ instruction.

```
1  #include <math.h>
2
3  /*
4  * 64 single precision registers      = s0 to s63
5  *                                   = d0 to d31
6  * This is twice that needed by my test hardware (VFPv3-D16).
7  */
8  #define NUM_SINGLE_PREC_REGS 64
9  int32_t fpu_registers[NUM_SINGLE_PREC_REGS];
10
11 /*
12 * Converts a single precision register number,
13 * such as '5' from the register 'r5', to a pointer into memory
14 * of where the value for that register is stored.
15 */
16 int32_t* sreg_to_bank_ptr(arm_reg reg) {
17     assert(ARM_REG_S0 <= reg && reg <= ARM_REG_S31);
18     return (int32_t*)fpu_registers + (reg - ARM_REG_S0);
19 }
20
21 // Sets a emulated floating-point register to the given 32-bit value
22 void set_sreg(arm_reg reg, int32_t val) {
23     *sreg_to_bank_ptr(reg) = val;
24 }
25
26 // Returns a value stored in an emulated floating-point register
27 int32_t get_sreg(arm_reg reg) {
28     return *sreg_to_bank_ptr(reg);
29 }
30
31 // Initialise emulator by setting the emulated registers to zero.
32 void emulator_init() {
33     memset(fpu_registers, 0, NUM_SINGLE_PREC_REGS * sizeof(int32_t));
34 }
35
36 /*
37 * Example of an emulation routine called by a trampoline.
38 * This is the emulation routine for the 'vadd.f32' instruction and will
```

```

39  * be called from a 'vadd.f32' trampoline. No C code calls this method
    as
40  * the machine code in the trampolines are generated at run-time.
41  */
42  void vadd_f32(int32_t Sd, int32_t Sn, int32_t Sm) {
43      float a, b;
44
45      /*
46       * Copy register values into local variables
47       * which are either on the stack or in scratch registers
48       * that are restored by the trampoline after this method returns.
49       */
50      int32_t Sn_val = get_sreg(Sn);
51      int32_t Sm_val = get_sreg(Sm);
52      memcpy(&a, &Sn_val, sizeof(int32_t));
53      memcpy(&b, &Sm_val, sizeof(int32_t));
54
55      // Perform the addition
56      float c = a + b;
57      int32_t result;
58
59      // Store the result back in a register
60      memcpy(&result, &c, sizeof(int32_t));
61      set_sreg(Sd, result);
62
63      printfdbg("vadd.f32 Sd:%d Sn:%d Sm:%d: %f + %f = %f\n", Sd, Sn, Sm,
64               a, b, c);
65      return;
66  }

```

4 rmaps.h

At the very beginning of instrumentation a file is opened to retrieve information about the memory layout. This opens, closes, and reads this file by interfacing with the ‘librunt’ library. Librunt’s callback here does the instrumentation for each line in the file.

```
1  #include <assert.h>
2  #include <stdio.h>
3  #include <string.h>
4  #include <stdint.h>
5  #include <link.h>
6  #include "maps.h"
7  #include "dso-meta.h"
8  #include "librunt.h"
9  #include "relf.h"
10
11  int MAPS_MAX_NUMLINEBUFS = 100;
12  int MAPS_BUF_SIZE = 1000;
13  FILE* maps_file;
14
15  /*
16   * In full transparency, this method resembles that of one of my
17   * supervisor's Dr. Stephen Kell.
18   * His project "libsystap" uses his other project "librunt" which is
19   * also a dependency for this project
20   * Due to the rigid intended use of the librunt functions "
21   * get_a_line_from_maps_fd"
22   * and "process_one_maps_entry", it was difficult to find a less similar
23   * way to implement this.
24   * Here is a link to that method "trap_all_mappings":
25   * https://github.com/stephenrkell/libsysstrap/blob/790
26   * cf958157520ce44afab0bcc2b0fcda9d168fe/example/trace-syscalls.c#L76
27   */
28  static void process_all_lines(int fd, void* callback) {
29      typedef char* linebuf_t;
30      linebuf_t linebufs[MAPS_MAX_NUMLINEBUFS];
31
32      /* Allocate line buffers and
33       * Read lines from "/proc/<self>/maps" into the line buffers
34       */
35      int line_counter = 0;
36      int newline_pos;
37      while (1) {
```

```

33     assert(line_counter < MAPS_MAX_NUMLINEBUFS);
34
35     linebufs[line_counter] = malloc(MAPS_BUF_SIZE);
36     newline_pos = get_a_line_from_maps_fd(linebufs[line_counter],
37     MAPS_BUF_SIZE, fd);
38     if (newline_pos == -1) break;
39     linebufs[line_counter][newline_pos] = '\0';
40
41     line_counter++;
42 }
43
44 int num_lines_read = line_counter; // to be explicit.
45
46 /* Process line buffers (including instrumenting relevant regions)
47 */
48 struct maps_entry mline;
49 int num_entries_skipped = 0;
50 int was_skipped;
51 for (int i = 0; i < num_lines_read; i++) {
52     if (0 == strncmp(linebufs[i], "00 ", 3)) { // if the region is
53         not memory mapped
54         was_skipped = 1;
55         printfdbg("Maps entry began with \"00 \" meaning it was
56         unmapped. Skipping.\n");
57     } else {
58         was_skipped = process_one_maps_entry(linebufs[i], &mline,
59         callback, NULL);
60     }
61     if (was_skipped) num_entries_skipped++;
62 }
63
64 /* Clean up */
65 for (int i = 0; i < num_lines_read; i++) {
66     free(linebufs[i]);
67 }
68
69 static int open_maps() {
70     maps_file = fopen("/proc/self/maps", "r");
71     int fd;
72     if ( maps_file == 0 || (fd = fileno(maps_file)) <= 2 ) {
73         printfdbg("Unable to open /proc/self/maps: Invalid file
74         descriptor.\n");
75         exit(1);
76     }
77     printfdbg("/proc/self/maps file descriptor: %d\n", fd);

```

```
73     return fd;
74 }
75
76 static void close_maps() {
77     fclose (maps_file);
78 }
```

5 debug-print.h

Some methods to print useful information when debugging.

```
1  /*
2  * Uncommenting/commenting this 'define' statement will
3  * enable or disable detailed console output useful for debugging.
4  */
5  // #define DO_DBG_PRINT
6
7  #ifdef DO_DBG_PRINT
8      #define printfdbg(...) printf (__VA_ARGS__)
9  #else
10     #define printfdbg(...)
11 #endif
12
13 extern long int etext;
14 extern long int edata;
15
16 // Returns 0 or 1 depending on whether an address is in the range [from
17 // , to]
18 static int is_in_range(long unsigned int from, long unsigned int to,
19                        long unsigned int addr) {
20     return ((from <= addr) && (addr <= to));
21 }
22
23 // Prints if a memory range contains the .text or .data sections
24 static int print_if_interesting_addr(long unsigned int from, long
25 unsigned int to) {
26     if (is_in_range(from, to, &etext)) {
27         printfdbg("Range contains .text\n");
28     } else if (is_in_range(from, to, &edata)) {
29         printfdbg("Range contains .data\n");
30     }
31     return 0;
32 }
33
34 // Print human-readable information about a programs section headers
35 static void print_section_headers(struct file_metadata* meta, ElfW(Shdr
36)* shdrs, int count) {
37     printfdbg("Section headers: \n");
38     for (int i = 0; i < count; i++) {
39         ElfW(Shdr)* shdr = shdrs + i;
40         if (shdr == NULL) {
```

```

37         printfdbg("\tshdrs[%d] does not exist. %d shdrs reported.\n",
38                 i, count);
39         break;
40     }
41     int shstrtab_index = shdr->sh_name;
42     char* section_name = meta->shstrtab + shstrtab_index;
43
44     char short_name[100];
45     if (strlen(section_name) == 0) strcpy(short_name, "(empty
46         header name)");
47     else short_name[0] = '\0';
48
49     if (shdr->sh_addr == 0) {
50         printfdbg("\t%s%s @ not loaded or base+0\n", short_name,
51             section_name);
52     } else {
53         printfdbg("\t%s%s \t@ base+%lx\n", short_name, section_name
54             , shdr->sh_addr);
55     }
56 }
57 return;
58 }
59
60 // Print human-readable information about a programs program headers
61 static void print_program_headers(struct file_metadata* meta, ElfW(Phdr)
62     * phdrs, int count) {
63     printfdbg("Program headers: \n");
64     for (int i = 0; i < count; i++) {
65         ElfW(Phdr)* phdr = phdrs + i;
66         if (phdr == NULL) {
67             printfdbg("phdrs[%d] does not exist. %d phdrs reported.\n",
68                 i, count);
69             break;
70         }
71         printfdbg("\t");
72         switch (phdr->p_type) {
73             case PT_NULL:
74                 printfdbg("PT_NULL - ignore the segment");
75                 break;
76             case PT_LOAD:
77                 printfdbg("PT_LOAD - loadable segment");
78                 break;
79             case PT_DYNAMIC:
80                 printfdbg("PT_DYNAMIC - dynamic linking info");
81                 break;
82             case PT_INTERP:

```



```

117         printfdbg("Error: could not retrieve metadata for file %s\n",
118                 ent->rest);
119         return 1;
120     }
121     ElfW(Ehdr)* ehdr = meta->ehdr;
122     printfdbg("%lx to %lx\n", ent->first, ent->second);
123     print_if_interesting_addr(ent->first, ent->second);
124     printfdbg("Privileges: %c%c%c%c\n", ent->r, ent->w, ent->x, ent->p)
125     ;
126     printfdbg("Inode: %d\n", ent->inode);
127     printfdbg("Filename: %s\n", meta->filename);
128     printfdbg("Magic: %s\n", ehdr->e_ident);
129     printfdbg("Arch: ");
130     switch(ehdr->e_machine) {
131     case EM_X86_64:
132         printfdbg("EM_X86_64");
133         break;
134     case EM_IA_64:
135         printfdbg("EM_IA_64");
136         break;
137     case EM_ARM:
138         printfdbg("EM_ARM");
139         break;
140     case EM_NONE:
141         printfdbg("EM_NONE");
142         break;
143     case EM_386:
144         printfdbg("EM_386");
145         break;
146     default: printfdbg("Unknown - %d", ehdr->e_machine);
147     }
148     printfdbg("\n");
149     printfdbg("Section header table file offset: %ld\n", ehdr->e_shoff)
150     ;
151     printfdbg("\tentry size: %d\n", ehdr->e_shentsize);
152     printfdbg("\tentry count: %d\n", ehdr->e_shnum);
153     printfdbg("Section header string table index: %d\n", ehdr->
154             e_shstrndx);
155     printfdbg("Segment header table file offset: %ld\n", ehdr->e_phoff)
156     ;
157     printfdbg("\tentry size: %d\n", ehdr->e_phentsize);
158     printfdbg("\tentry count: %d\n", ehdr->e_phnum);
159     print_section_headers(meta, meta->shdrs, ehdr->e_shnum);

```

```
158     print_program_headers(meta, meta->phdrs, ehdr->e_phnum);
159
160     printfdbg("\n");
161     return 0;
162 }
```